

Jin-Soo Kim
(jinsoo.kim@snu.ac.kr)

Systems Software &
Architecture Lab.

Seoul National University

Spring 2020

Solid State Drives (SSDs)




Memory Types

FLASH

- High-density
- Low-cost
- High-speed
- Low-power
- High reliability

DRAM

- High-density
- Low-cost
- High-speed
- High-power



EPROM

- Non-volatile
- High-density
- Ultraviolet light for erasure

EEPROM

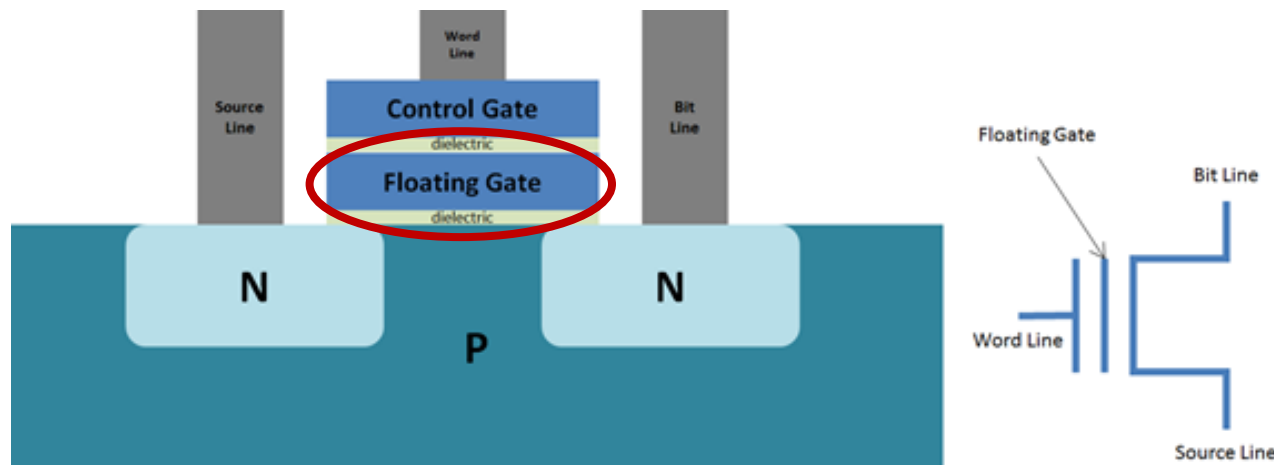
- Non-volatile
- Lower reliability
- Higher cost
- Lowest density
- Electrically byte-erasable

ROM

- High-density
- Reliable
- Low-cost
- Suitable for high production with stable code

Flash Memory Cell

- Transistor with floating gate
 - The floating gate is insulated all around with an oxide layer
 - Electrons trapped in the floating gate can remain for up to years





Flash Memory Characteristics

■ Erase-before-write

- Read
- Write or Program: $1 \rightarrow 0$
- Erase: $0 \rightarrow 1$

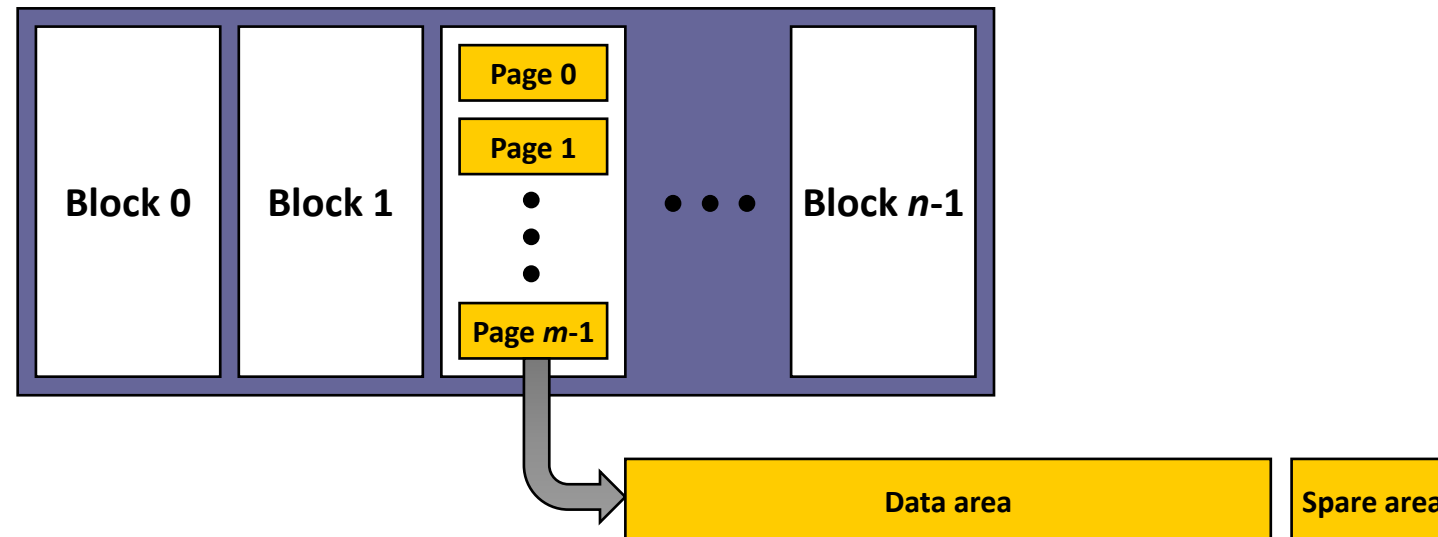
■ Bulk erase

- Program unit:
 - NOR: byte or word
 - NAND: page
- Erase unit: _____




Logical View of NAND Flash

- A collection of **blocks**
- Each block has a number of **pages**
- The size of a block or a page depends on the technology (but, it's getting larger)



NAND Flash Types

- SLC NAND
 - Single Level Cell
 - 1 bit/cell
- MLC NAND
 - Multi Level Cell (misnomer)
 - 2 bits/cell
- TLC NAND
 - Triple Level Cell
 - 3 bits/cell
- 3D NAND




NAND Applications

- Universal Flash Drives (UFDs)
- Flash cards
 - CompactFlash, MMC, SD, Memory stick, ...
- Smartphones
 - eMMC (Embedded MMC)
 - UFS (Universal Flash Storage)
- SSDs (Solid State Drives)
- Other embedded devices
 - MP3 players, Digital TVs, Set-top boxes, Car navigators, ...



Anatomy of an SSD

- Samsung 850 Evo



<http://www.anandtech.com/show/9451/the-2tb-samsung-850-pro-evo-ssd-review>

HDDs vs. SSDs

Feature	SSD (Samsung)	HDD (Seagate)
Model	MZ-75E2T0B (850 Evo)	ST2000LM003 (SpinPoint M9T)
Capacity	2TB (128Gb 32-Layer 3D V-NAND TLC x 16 die/channel x 8 channels)	2TB (3 Discs, 6 Heads, 5400 RPM)
Form factor	2.5", 66g	2.5", 130g
DRAM	2 GB	32 MB
Host interface	SATA-3 (6.0 Gbps)	SATA-3 (6.0 Gbps)
Power consumption (Active / Idle / Sleep)	3.7, 4.7 W / 0.5 W / 0.05 W	2.3 W / 0.7 W / 0.18 W
Performance 850 Evo ¹ : Sequential: 128KB/QD2 Random: 4KB/QD32 M9T ² : Sequential: 2MB Random: 4KB	Sequential read: 544 MB/s Sequential write: 520 MB/s Random read: 97,687 IOPS Random write: 89,049 IOPS Random read: 11,335 IOPS (QD1) Random write: 38,433 IOPS (QD1)	Sequential read: 124 MB/s Sequential write: 124 MB/s Random read: 56 IOPS Random write: 98 IOPS Power-on to ready: 3.5 sec Average seek: 12/14 ms Average latency: 5.6 ms
Price ³	940,910 won (470won/GB)	175,900 won (88won/GB)

¹ <http://www.tomshardware.com/reviews/samsung-850-evo-850-pro-2tb-ssd,4205.html>

² http://www.storagereview.com/samsung_spinpoint_m9t_hard_drive_review ³ <http://www.enuri.com> (As of May. 26, 2020)

State-of-the-Art @ 2018

삼성전자, 세계 최초 '30.72TB SAS SSD' 양산

2018/02/20

공유하기



- 2.5"
- 1TB V-NAND x32
- 4GB DRAM x 10
- Sequential read: 2100MB/s
- Sequential write: 1700MB/s

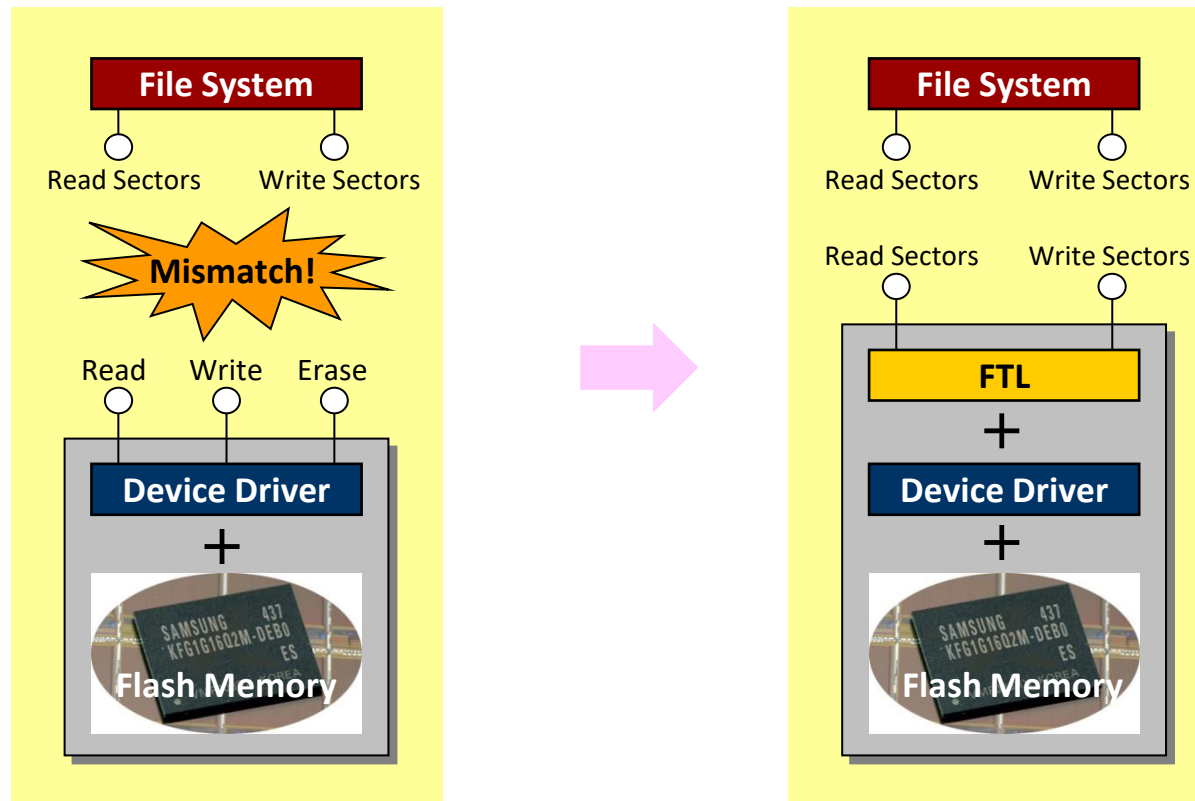


NAND Constraints


- No in-place update
 - Require sector remapping (or address translation)
- Bit errors
 - Require the use of error correction codes (ECCs)
- Bad blocks
 - Factory-marked and run-time bad blocks
 - Require bad block remapping
- Limited program/erase cycles
 - $< 100K$ for SLCs, $< 3K$ for MLCs, $< 1K$ for TLCs
 - Require wear-leveling

Flash Translation Layer (FTL)

- A software layer to make NAND flash fully emulate traditional block devices (e.g., disks)

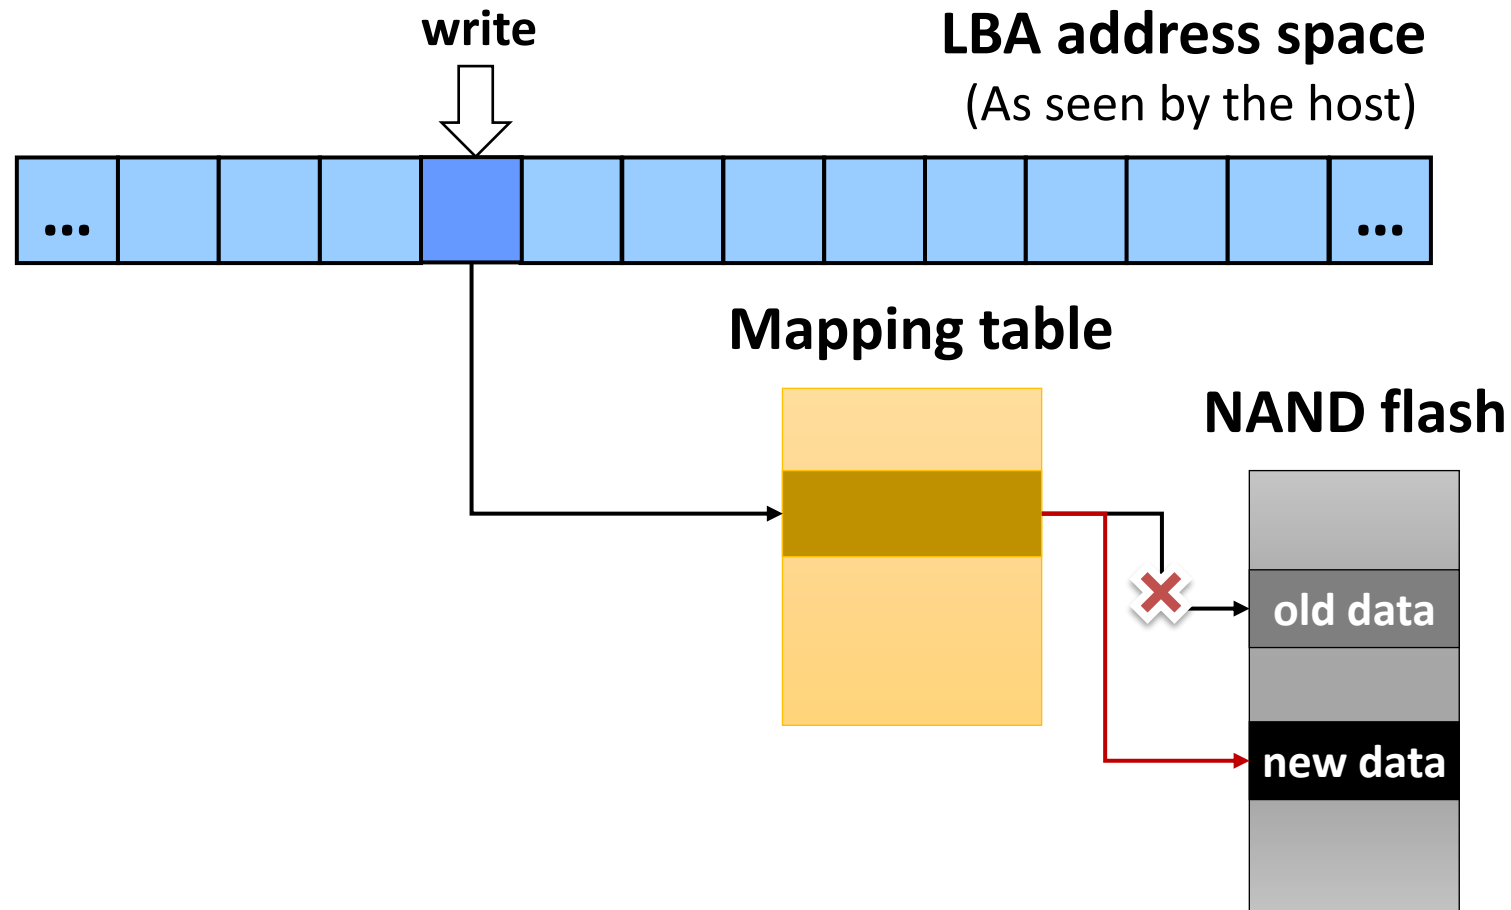


SSD Internals



Address Mapping

- Required since flash pages cannot be overwritten



Example: Page Mapping

- Flash configuration

- Page size: 4KB
- # of pages / block = 4


- Current state

- Written to page 0, 1, 2, 8, 4, 5

- Reading page 5

Logical page #5

0000000101



Example: Page Mapping

- Flash configuration
 - Page size: 4KB
 - # of pages / block = 4
- Current state
 - Written to page 0, 1, 2, 8, 4, 5
- New requests (in order)
 - Write to page 9
 - Write to page 3
 - Write to page 5

Page Map Table		Data Block		PPN
0	0	PBN: 0	0	0
1	1		1	1
2	2		2	2
3			8	3
4	4	PBN: 1	4	4
5	5		5	5
6				6
7				7
8	3	PBN: 2		8
9				9
10				10
11				11
		PBN: 3		12
				13
				14
				15

Example: Page Mapping

- Flash configuration
 - Page size: 4KB
 - # of pages / block = 4
- Current state
 - Written to page 0, 1, 2, 8, 4, 5
- New requests (in order)
 - Write to page 9
 - Write to page 3
 - Write to page 5

Page Map Table		Data Block		PPN
0	0	PBN: 0	0	0
1	1		1	1
2	2		2	2
3			8	3
4	4	PBN: 1	4	4
5	5		5	5
6			9	6
7				7
8	3	PBN: 2		8
9	6			9
10				10
11				11
		PBN: 3		12
				13
				14
				15

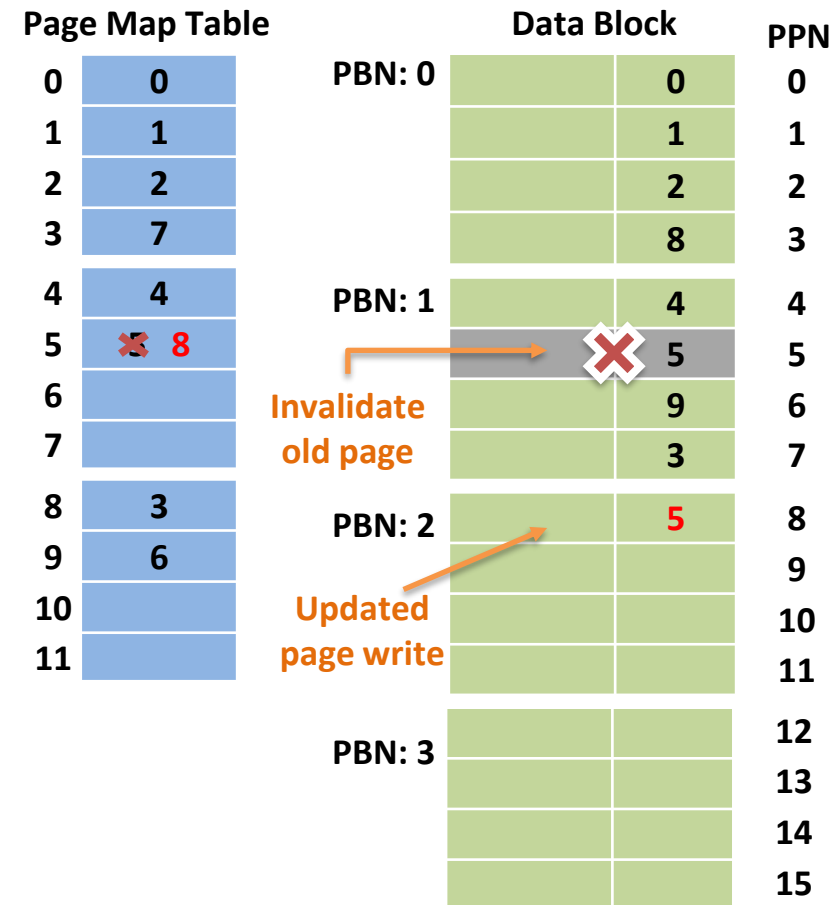
Example: Page Mapping

- Flash configuration
 - Page size: 4KB
 - # of pages / block = 4
- Current state
 - Written to page 0, 1, 2, 8, 4, 5
- New requests (in order)
 - Write to page 9
 - Write to page 3
 - Write to page 5

Page Map Table		Data Block		PPN
0	0	PBN: 0	0	0
1	1		1	1
2	2		2	2
3	7		8	3
4	4	PBN: 1	4	4
5	5		5	5
6			9	6
7			3	7
8	3	PBN: 2		8
9	6			9
10				10
11				11
		PBN: 3		12
				13
				14
				15

Example: Page Mapping

- Flash configuration
 - Page size: 4KB
 - # of pages / block = 4
- Current state
 - Written to page 0, 1, 2, 8, 4, 5
- New requests (in order)
 - Write to page 9
 - Write to page 3
 - Write to page 5**



Garbage Collection

- Garbage collection (GC)
 - Eventually, FTL will run out of blocks to write to
 - GC must be performed to reclaim free space
 - Actual GC procedure depends on the mapping scheme
- GC in page-mapping FTL
 - Select victim block(s)
 - Copy all valid pages of victim block(s) to free block
 - Erase victim block(s)
 - Note: At least one free block should be reserved for GC

Example: GC in Page Mapping

■ Current state

- Written to page 0, 1, 2, 8, 4, 5
- Written to page 9, 3, 5

■ New requests (in order)

- Write to page 8
- Write to page 9
- Write to page 3
- Write to page 1
- Write to page 4

Page Map Table			Data Block		PPN
0	0	PBN: 0		0	0
1	1			1	1
2	2			2	2
3	7			8	3
4	4	PBN: 1		4	4
5	8			5	5
6				9	6
7				3	7
8	3	PBN: 2		5	8
9	6				9
10					10
11					11
		PBN: 3			12
					13
					14
					15

Example: GC in Page Mapping

■ Current state

- Written to page 0, 1, 2, 8, 4, 5
- Written to page 9, 3, 5

■ New requests (in order)

- Write to page 8
- Write to page 9
- Write to page 3
- Write to page 1
- Write to page 4

Page Map Table		Data Block		PPN
0	0	PBN: 0	0	0
1	1		1	1
2	2		2	2
3	7		8	3
4	4	PBN: 1	4	4
5	8		5	5
6			9	6
7			3	7
8	9	PBN: 2	5	8
9	6		8	9
10				10
11				11
		PBN: 3		12
				13
				14
				15

Example: GC in Page Mapping

■ Current state

- Written to page 0, 1, 2, 8, 4, 5
- Written to page 9, 3, 5

■ New requests (in order)

- Write to page 8
- **Write to page 9**
- Write to page 3
- Write to page 1
- Write to page 4

Page Map Table		Data Block		PPN
0	0	PBN: 0	0	0
1	1		1	1
2	2		2	2
3	7		8	3
4	4	PBN: 1	4	4
5	8		5	5
6			9	6
7			3	7
8	9	PBN: 2	5	8
9	10		8	9
10			9	10
11				11
		PBN: 3		12
				13
				14
				15

Example: GC in Page Mapping

■ Current state

- Written to page 0, 1, 2, 8, 4, 5
- Written to page 9, 3, 5

■ New requests (in order)

- Write to page 8
- Write to page 9
- **Write to page 3**
- Write to page 1
- Write to page 4

Page Map Table		Data Block		PPN
0	0	PBN: 0	0	0
1	1		1	1
2	2		2	2
3	11		X 8	3
4	4	PBN: 1	4	4
5	8		X 5	5
6			X 9	6
7			X 3	7
8	9	PBN: 2	5	8
9	10		8	9
10			9	10
11			3	11
		PBN: 3		12
			Spare block	13
				14
				15

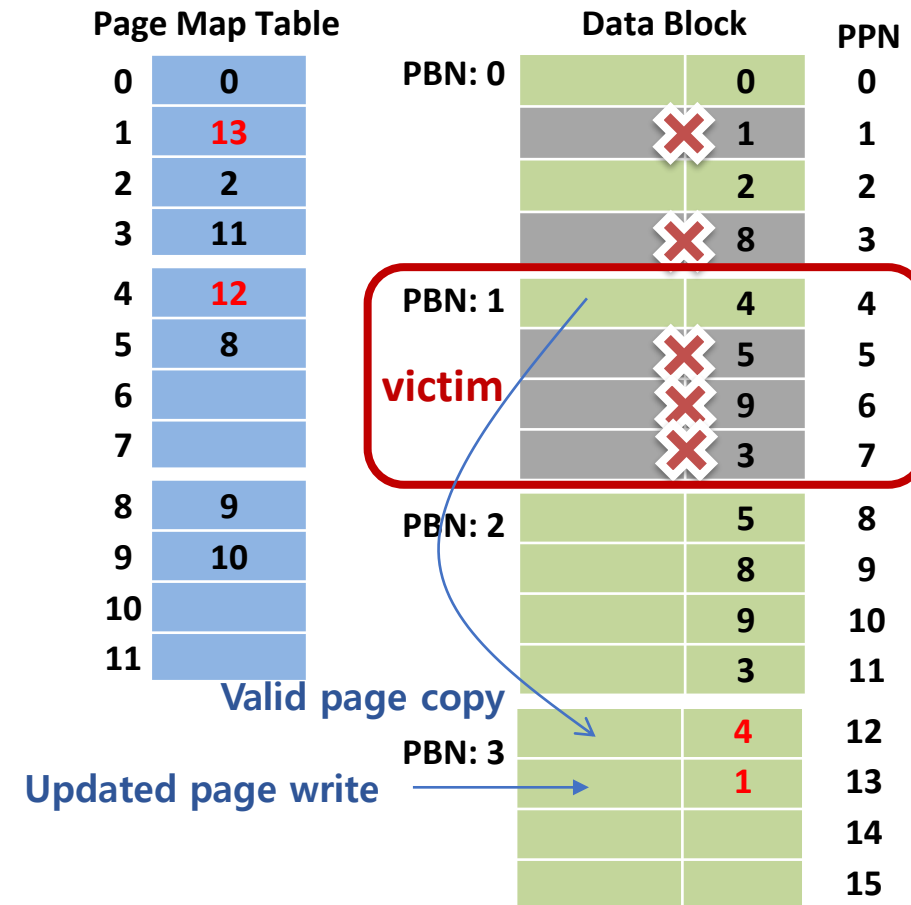
Example: GC in Page Mapping

■ Current state

- Written to page 0, 1, 2, 8, 4, 5
- Written to page 9, 3, 5

■ New requests (in order)

- Write to page 8
- Write to page 9
- Write to page 3
- **Write to page 1**
- Write to page 4



Example: GC in Page Mapping

■ Current state

- Written to page 0, 1, 2, 8, 4, 5
- Written to page 9, 3, 5

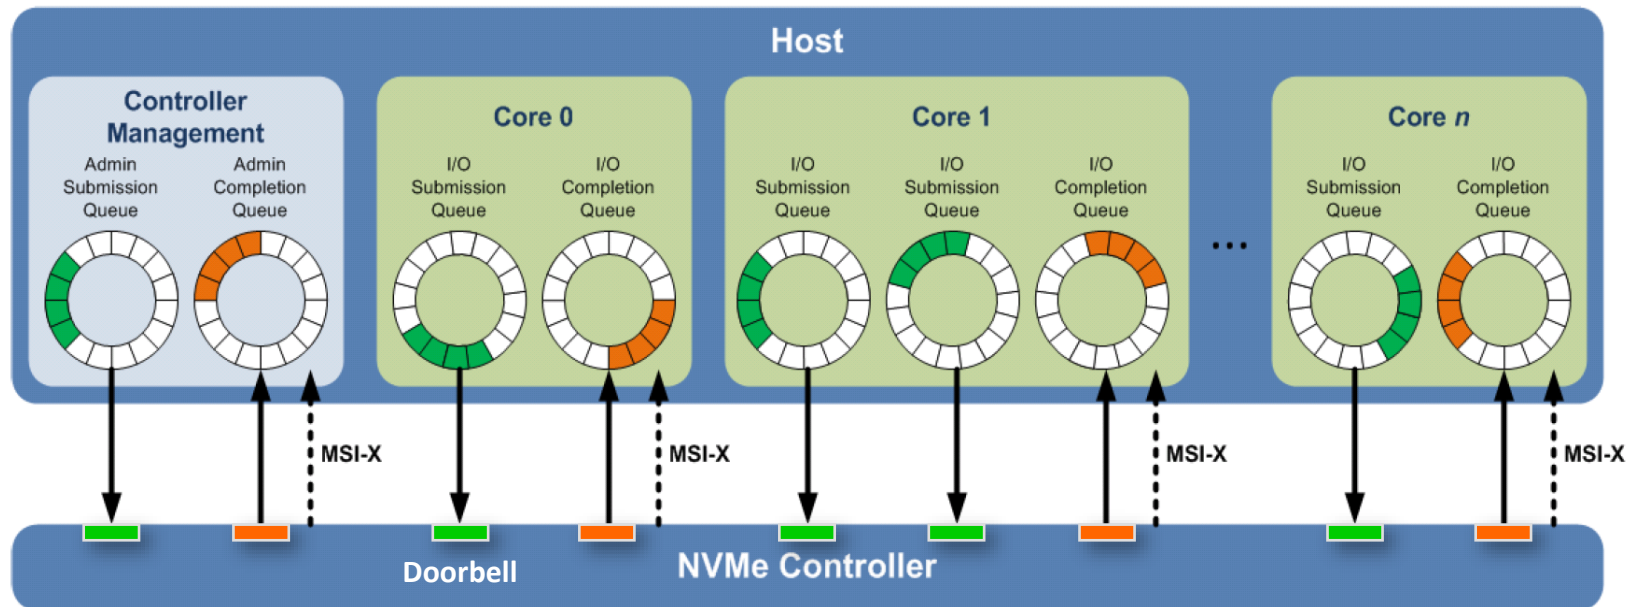
■ New requests (in order)

- Write to page 8
- Write to page 9
- Write to page 3
- Write to page 1
- **Write to page 4**

Page Map Table		Data Block		PPN
0	0	PBN: 0	0	0
1	13		1	1
2	2		2	2
3	11		8	3
4	14	PBN: 1		4
5	8			5
6				6
7				7
8	9	PBN: 2	5	8
9	10		8	9
10			9	10
11			3	11
		PBN: 3	4	12
			1	13
			4	14
				15

NVMe SSD

- PCIe-based (PCIe Gen. 3: 1GB/s per lane, up to 32 lanes)
- Deep queue: 64K commands per queue, up to 64K queues
- Streamlined command set: only 13 required commands
- One register write to issue a command (“doorbell”)



OS Implications

- NAND flash has different characteristics compared to disks
 - No seek time
 - Asymmetric read/write access times
 - No in-place-update
 - Good sequential read/write and random read performance, but bad random write performance
 - Wear-leveling
 - ...
- Traditional operating systems have been optimized for disks. What should be changed?

SSD Support in OS

- Turn off “defragmentation” for SSDs
- New “TRIM” command
 - Remove-on-delete
- Simpler I/O scheduler
- Align file system partition with SSD layout
- Flash-aware file systems (e.g., F2FS in Linux)
- Larger block size (4KB)
- New “multi-stream” interface
- ...

Beauty and the Beast

- NAND Flash memory is a beauty
 - Small, light-weight, robust, low-cost, low-power, non-volatile device
- NAND Flash memory is a beast
 - No in-place-update
 - Much slower program/erase operations
 - Erase unit > read/write unit
 - Bit errors
 - Limited lifetime etc.
- Software support is essential for performance and reliability!

